

Online Object Detection

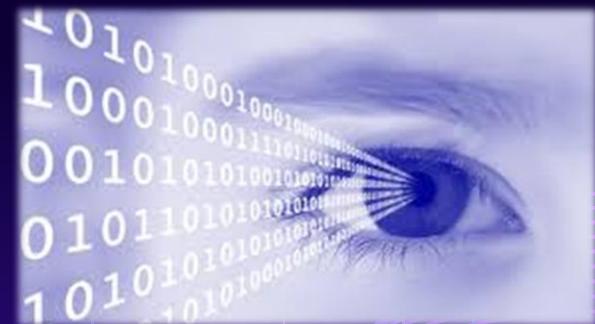
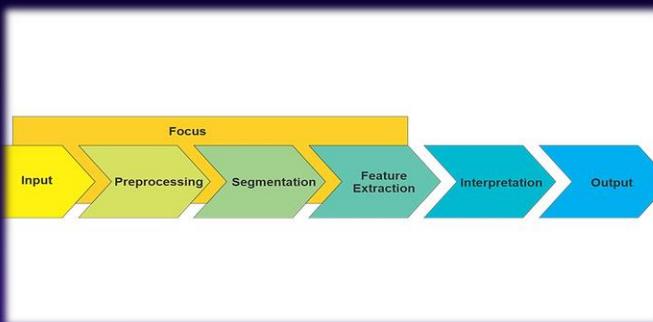
(Image Processing)

M.M.Mirzaei

Instructor:Ms.Kamali.Yazdi

It's all about digital image processing (DIP)

Digital Image Processing means processing digital image by means of a digital computer. We can also say that it is a use of computer algorithms, in order to get enhanced image either to extract some useful information.



Requires the knowledge of basic calculus, probability, and differential equations for mathematical calculation purposes. Also, there is a basic knowledge of programming like MATLAB, C++, etc.

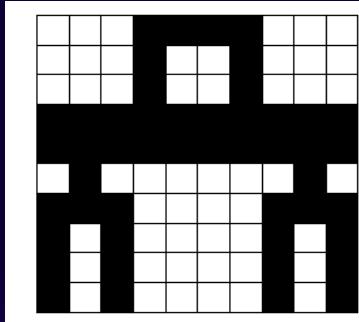
What is an Digital Image?

Projection of 3D scene in 2D plane.

An image is defined as a two-dimensional function, $F(x,y)$, where x and y are spatial coordinates, and the amplitude of F at any pair of coordinates (x,y) is called the **intensity** of that image at that point. When x,y , and amplitude values of F are finite, we call it a **digital image**.

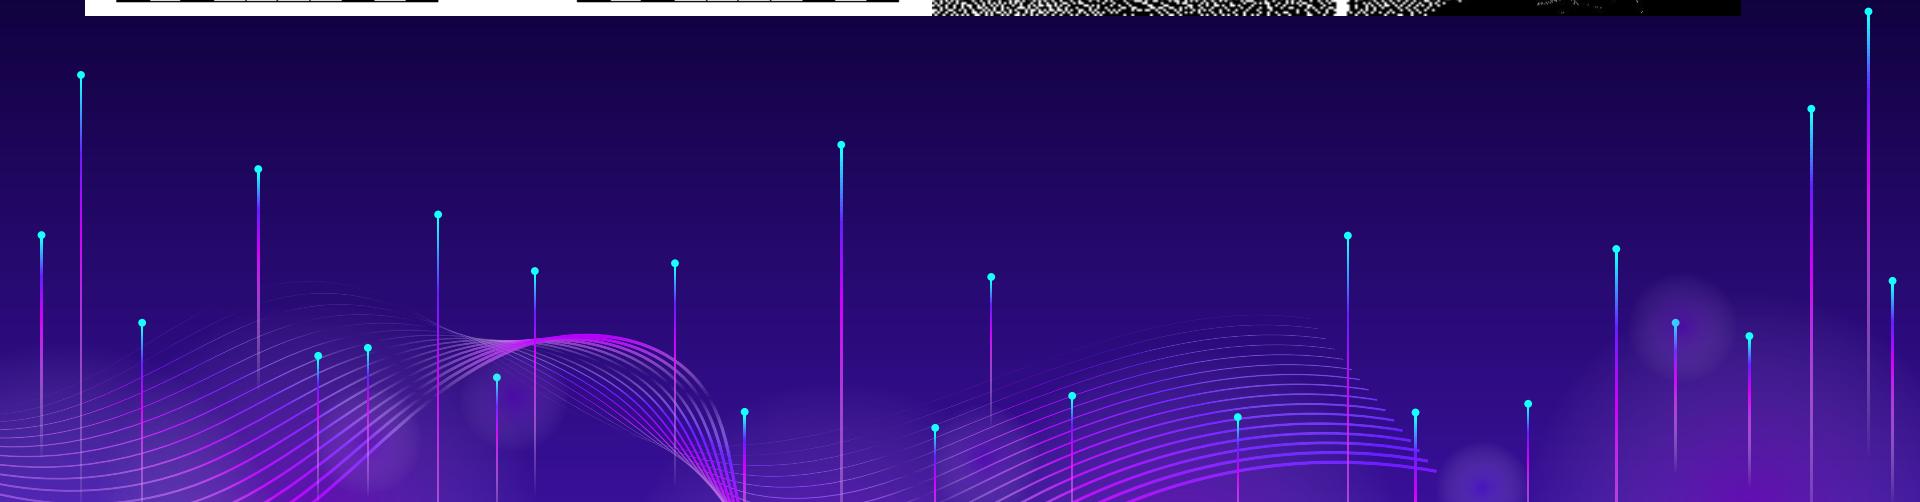
1. **BINARY IMAGE**– The binary image as its name suggests, contain only two pixel elements i.e 0 & 1,where 0 refers to black and 1 refers to white. This image is also known as Monochrome.
2. **BLACK AND WHITE IMAGE**– The image which consist of only black and white color is called BLACK AND WHITE IMAGE.
3. **8 bit COLOR FORMAT**– It is the most famous image format. It has 256 different shades of colors in it and commonly known as Grayscale Image. In this format, 0 stands for Black, and 255 stands for white, and 127 stands for gray.
4. **16 bit COLOR FORMAT**– It is a color image format. It has 65,536 different colors in it. It is also known as High Color Format. In this format the distribution of color is not as same as Grayscale image.
 - 4.1 RGB - A pixel is made up of 3 integers between 0 to 255 (the integers represent the intensity of red, green, and blue).
 - 4.2 RGBA - It is an extension of RGB with an added alpha field, which represents the opacity of the image.

BINARY IMAGE

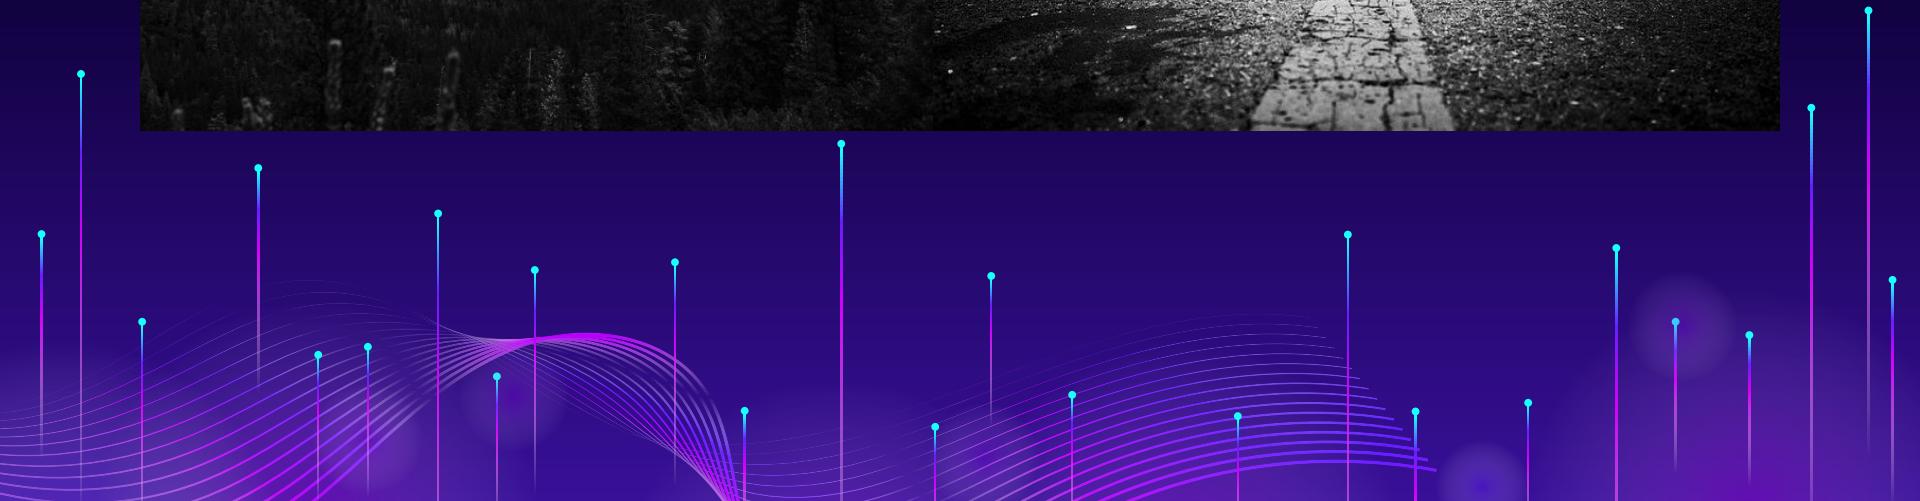


→

0	0	0	1	1	1	1	0	0	0
0	0	0	1	0	0	1	0	0	0
0	0	0	1	0	0	1	0	0	0
1	1	1	1	1	1	1	1	1	1
1	1	1	1	1	1	1	1	1	1
0	1	0	0	0	0	0	0	1	0
1	1	1	0	0	0	0	1	1	1
1	0	1	0	0	0	0	1	0	1
1	0	1	0	0	0	0	1	0	1
1	0	1	0	0	0	0	1	0	1



Black and White Image



8 bit Color Format

16 bit Color Format

Difference between
8-bit and 16 bit color format in DIP

æ

8-bit Color Representation



16-bit Color Representation

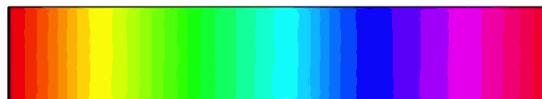


Image Processing Steps

- Image Acquisition
- Image Enhancement
- Image Restoration
- Color Image Processing
- Image Segmentation

The background features a dark purple gradient with a subtle wavy pattern. Overlaid on this are numerous thin, light blue lines that form a grid-like structure, with small cyan dots at their intersections, resembling a sensor array or a digital grid.

01

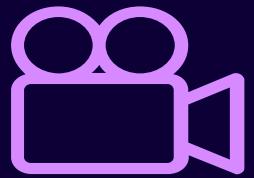
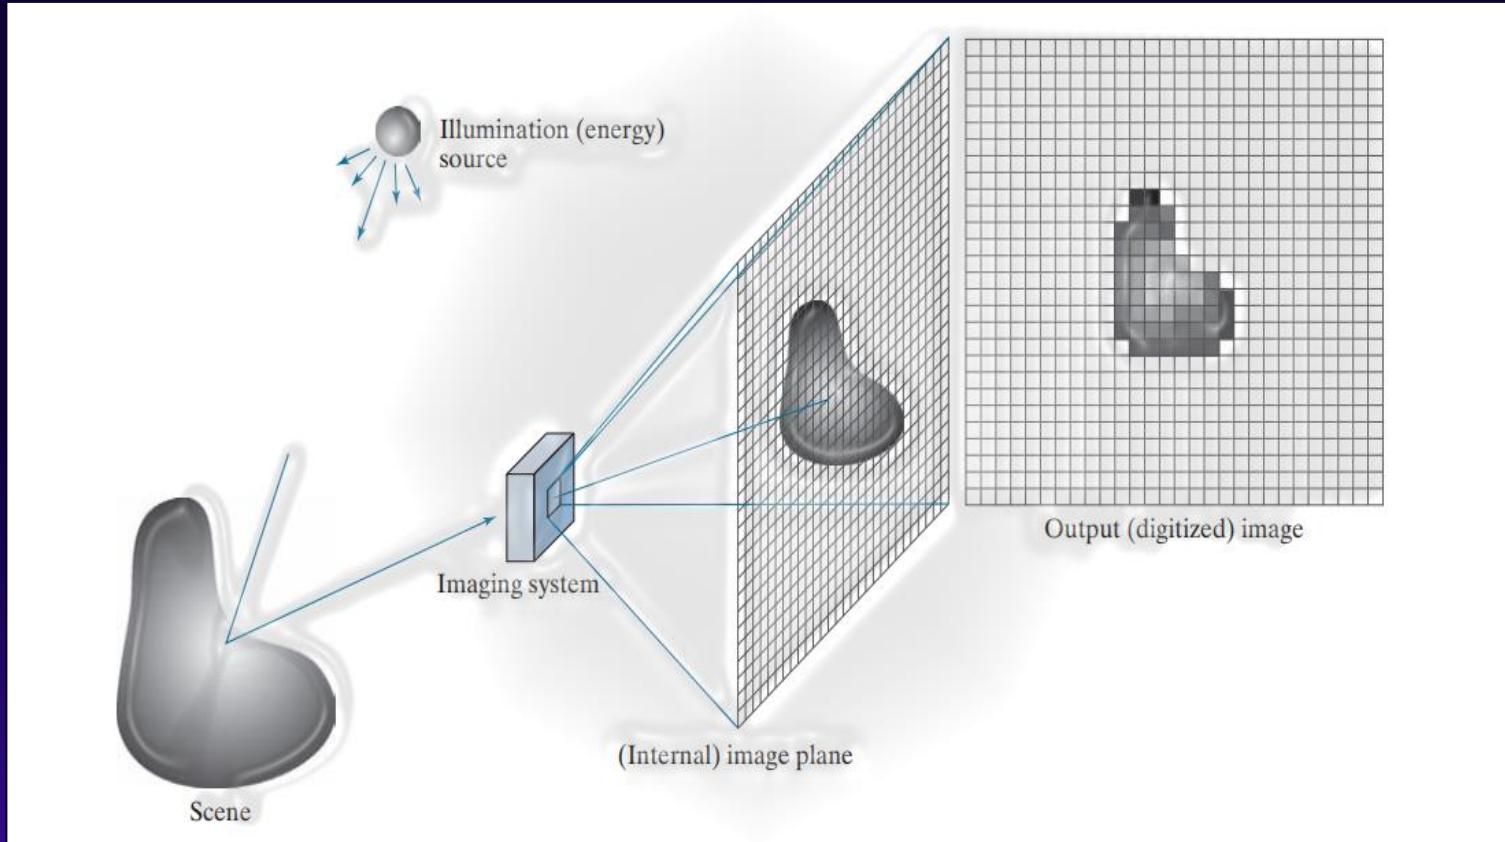
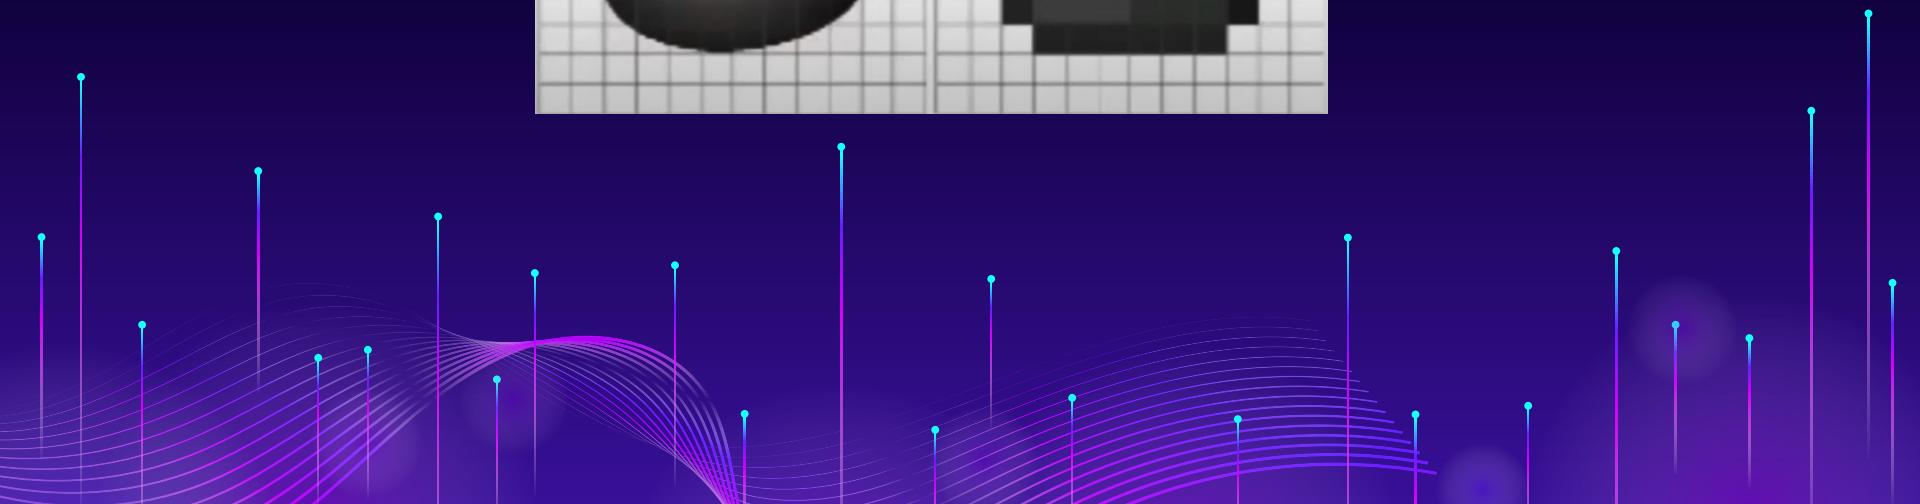
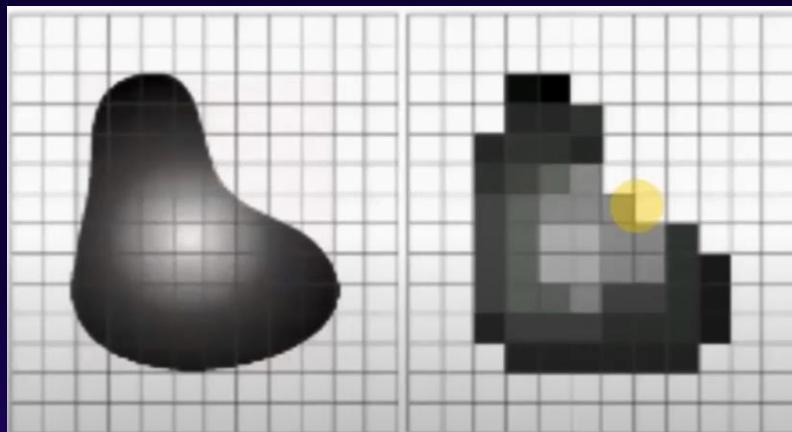


Image Acquisition

The process of capturing visual information from the real world and converting it into a digital image that can be processed by computers



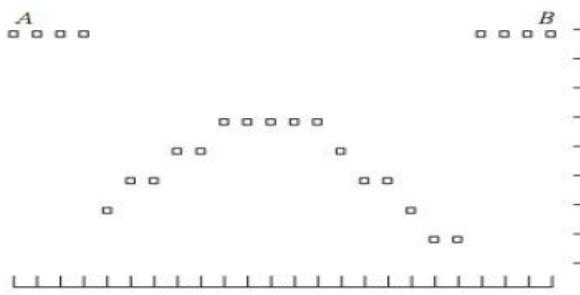
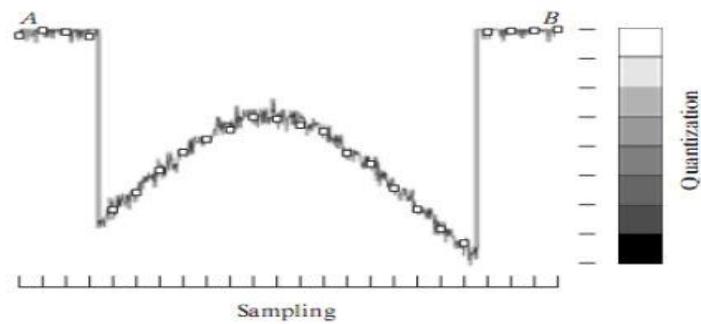
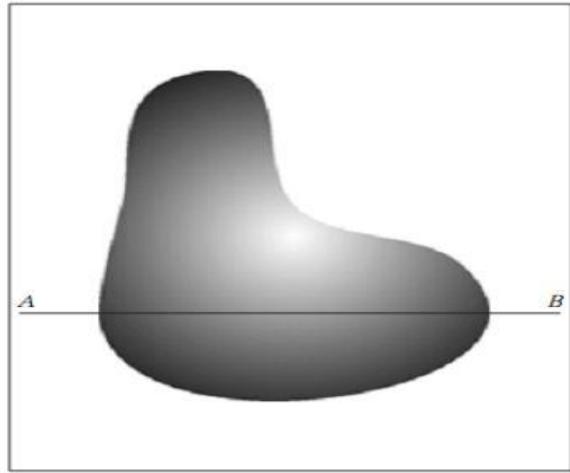
Convert Analog Image To Digital Image

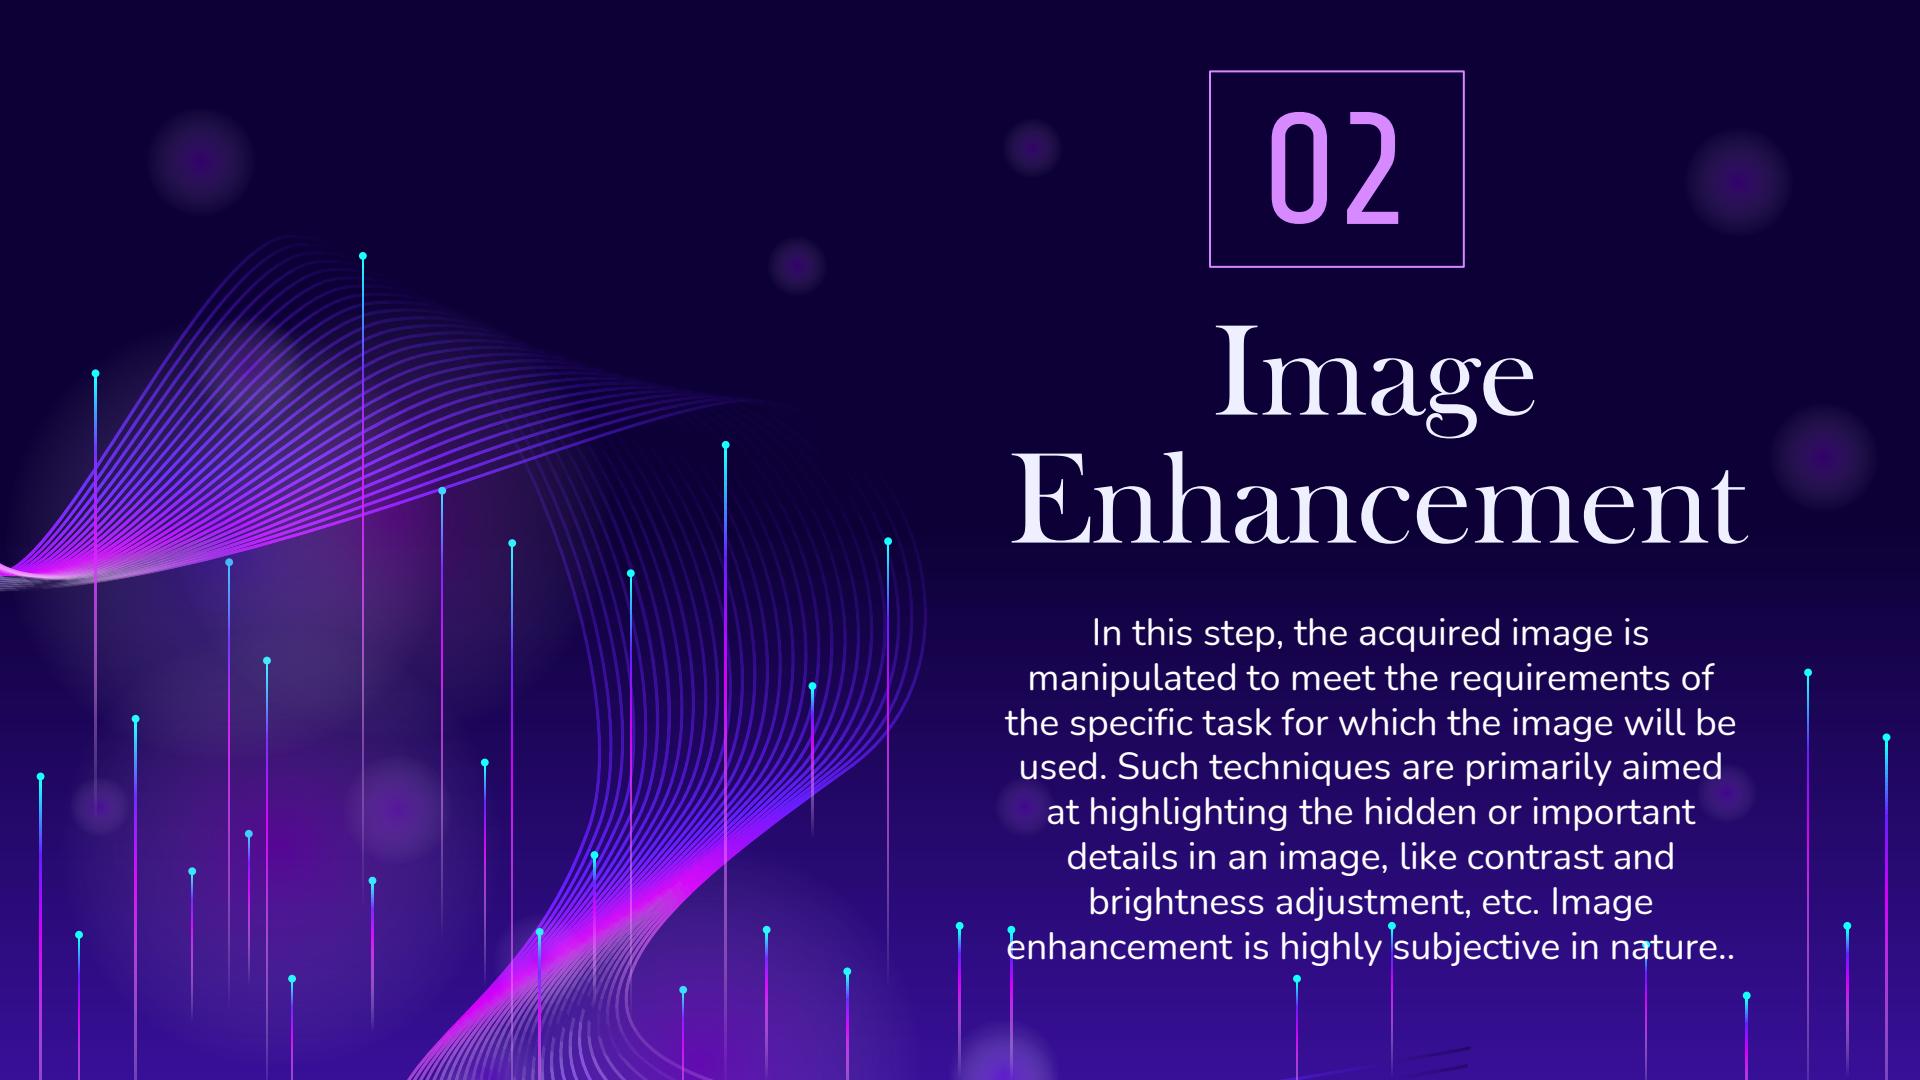


Digitizing an Image

Sampling

Quantization





02

Image Enhancement

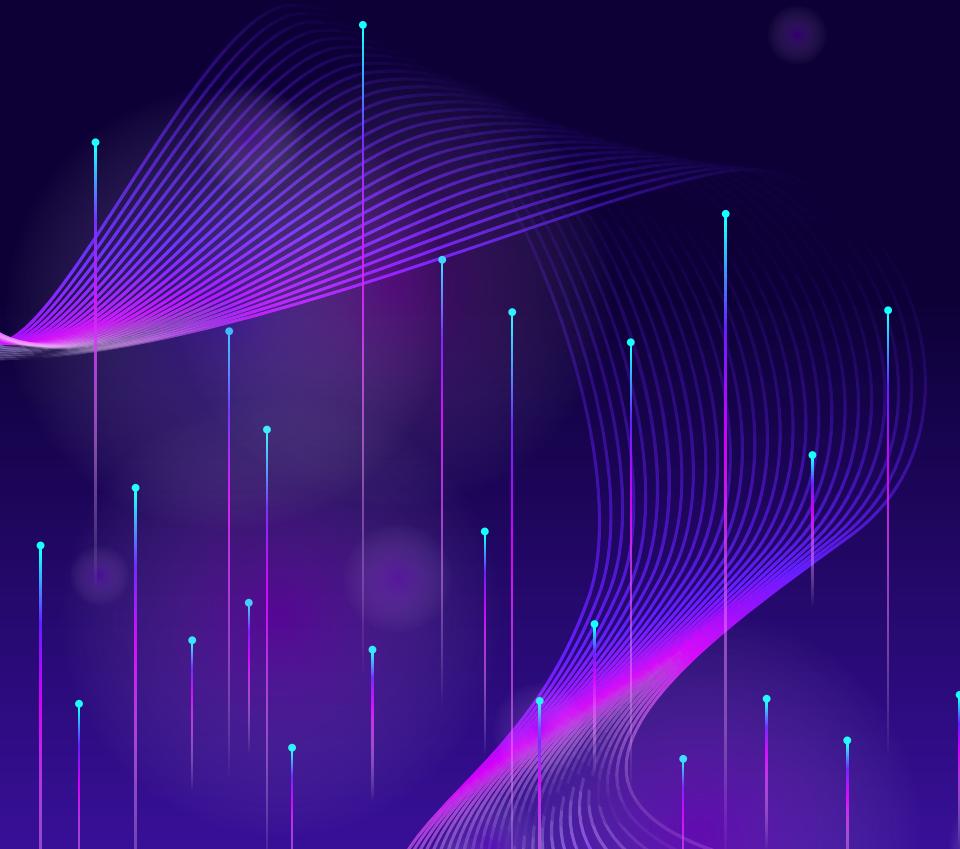
In this step, the acquired image is manipulated to meet the requirements of the specific task for which the image will be used. Such techniques are primarily aimed at highlighting the hidden or important details in an image, like contrast and brightness adjustment, etc. Image enhancement is highly subjective in nature..



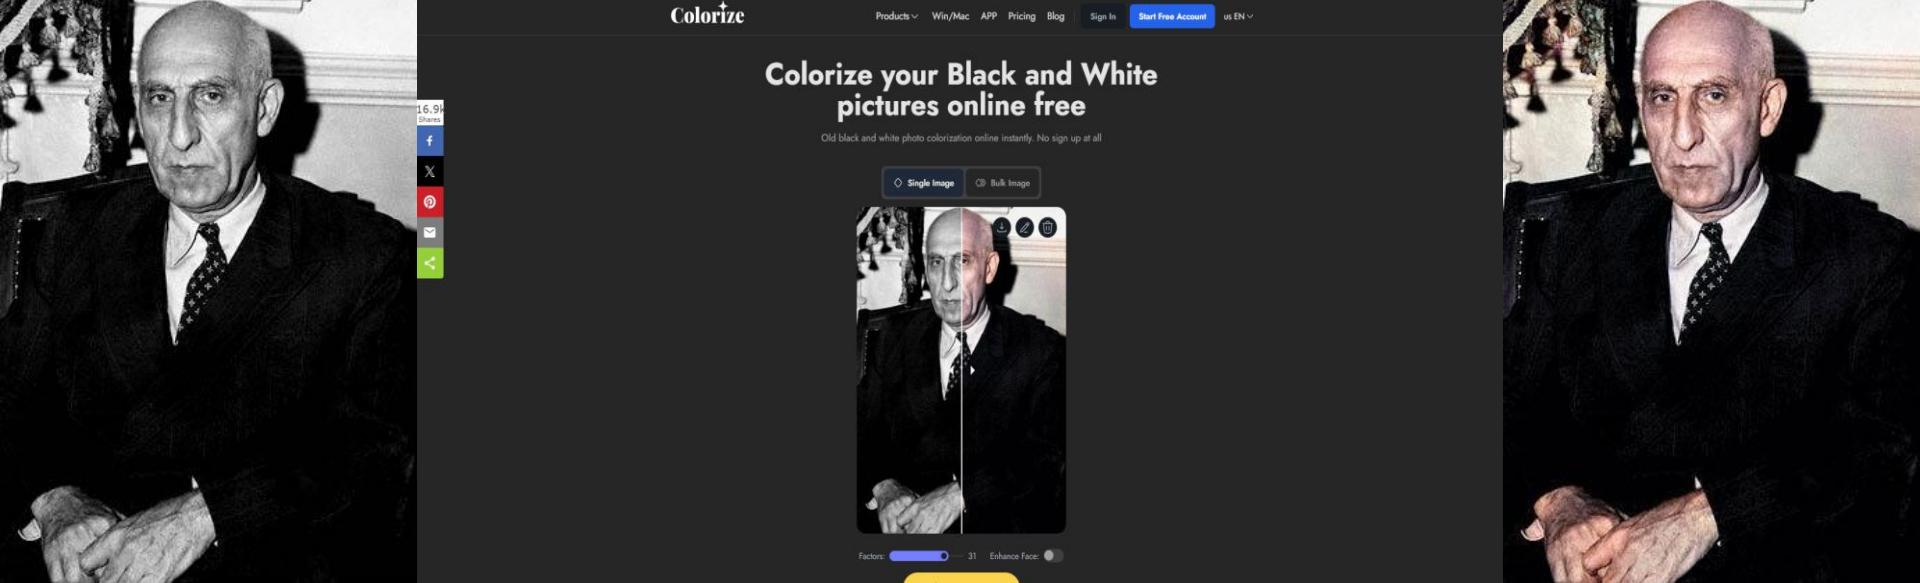
03

Image Restoration

Image restoration is the process of improving the appearance of an image. However, unlike image enhancement, image restoration is done using certain **mathematical** or **probabilistic** models.







<https://imagecolorizer.com/colorize>

04

Color Image Processing

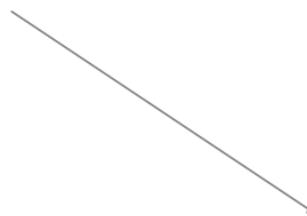
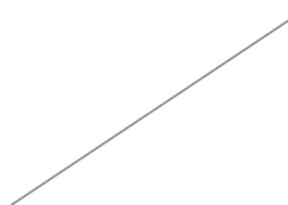
This step aims at handling the processing of colored images (16-bit RGB or RGBA images), for example, performing color correction or color modeling in images.



—

—

Blue				
Green	123	94	83	2
Red	123	94	83	4
123	94	83	2	92
34	44	187	92	4
34	76	232	124	4
67	83	194	202	





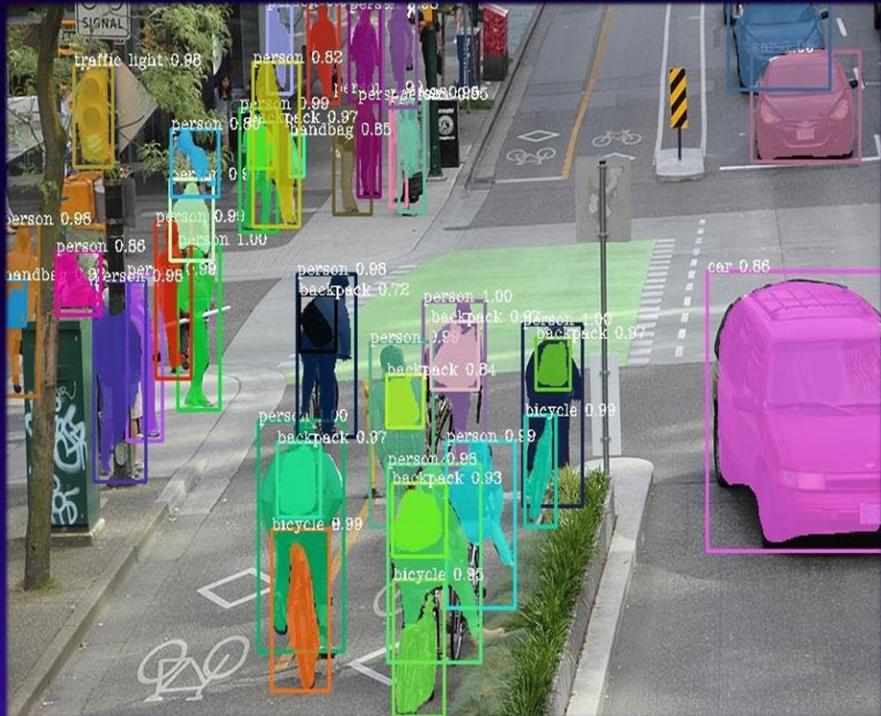
05

Image Segmentation

Divides an image into distinct groups of pixels, called segments, to help with object detection and other tasks. It's a key step in image analysis and processing, and is used in many fields, including medical imaging, autonomous driving, and background removal

Image Segmentation

- **Thresholding**
 - Clustering methods
 - Dual clustering method
 - Motion and interactive segmentation
 - Compression-based methods
 - Histogram-based methods
 - **Edge detection**
 - Isolated Point Detection
 - Region-growing methods
 - Variational methods
 - Graph partitioning methods



Object Recognition Methods

Object Recognition Using Machine Learning

- **HOG (Histogram of oriented Gradients) feature Extractor and SVM (Support Vector Machine) model**
Before the era of deep learning, it was a state-of-the-art method for object detection. It takes histogram descriptors of both **positive (images that contain objects)** and **negative (images that does not contain objects)**
- **Viola-Jones algorithm:** This algorithm is widely used for **face detection** in the image or **real-time**. It performs **Haar-like feature extraction** from the image. This generates a large number of features. These features are then passed into a boosting classifier.

Viola-Jones is that it has a detection time of 2 fps which can be used in a real-time face recognition system.



Object Recognition Using Deep Learning

Convolution Neural Network (CNN) is one of the most popular ways of doing object recognition. It is widely used and most state-of-the-art neural networks used this method for various object recognition related tasks such as image classification. This CNN network takes an image as input and outputs the probability of the different classes. If the object present in the image then it's output probability is high else the output probability of the rest of classes is either negligible or low. The advantage of Deep learning is that we don't need to do feature extraction from data as compared to machine learning.

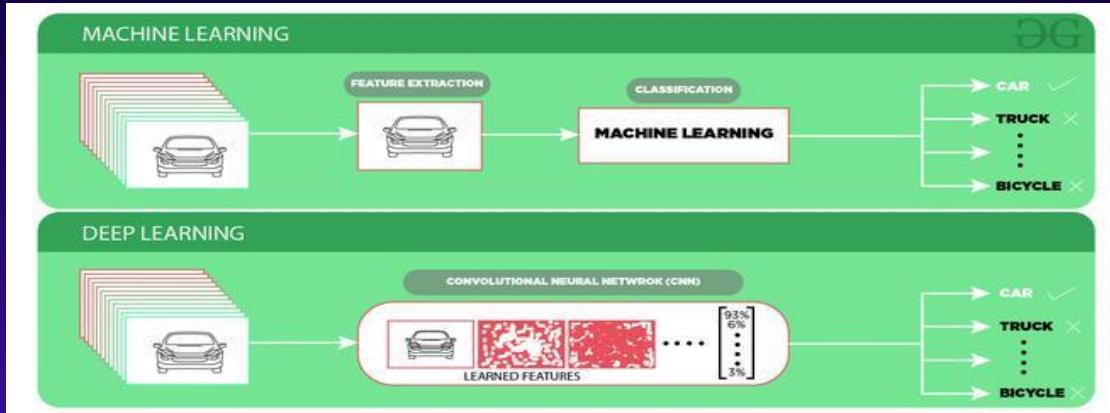


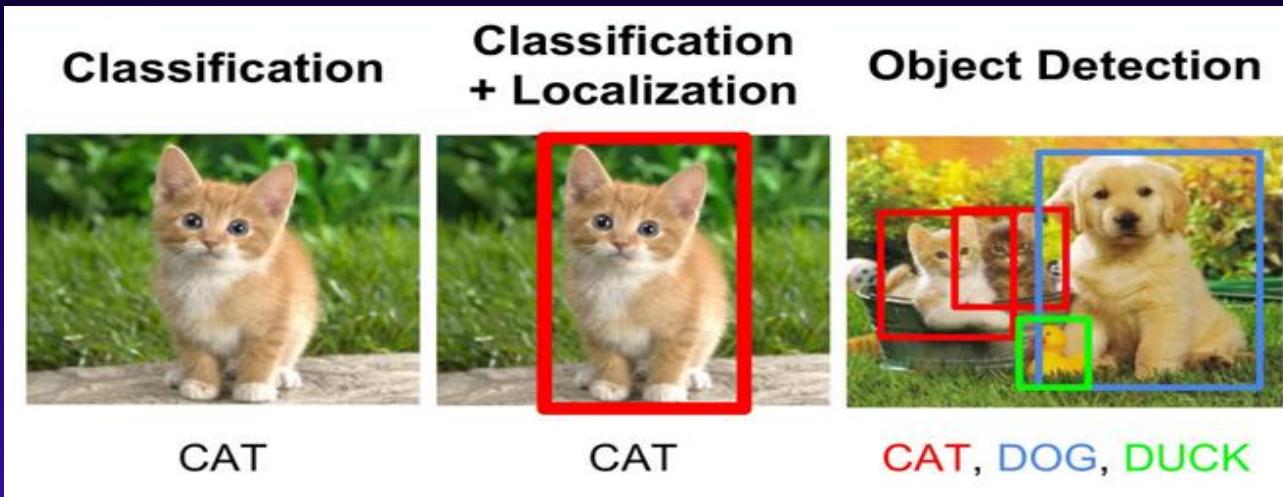
Image Classification

In Image classification, it takes an image as an input and outputs the classification **label** of that image with some metric (probability, loss, accuracy, etc). For Example: An image of a cat can be classified as a class label “cat” or an image of Dog can be classified as a class label “dog” with some probability.



Object Localization

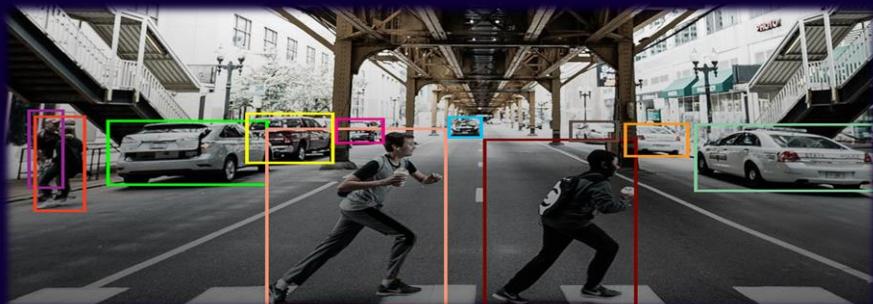
This algorithm locates the presence of an object in the image and represents it with a **bounding box**. It takes an image as input and outputs the location of the bounding box in the form of (position, height, and width).



What is an object detection?

Object detection is the process of taking a machine that has a camera on it and training it on one or more different objects afterwards that machine will be able to recognize the objects and report back the position of the object or objects within its field of view.

Object Detection algorithms act as a **combination of image classification and object localization**. It takes an image as input and produces one or more bounding boxes with the class label attached to each bounding box. These algorithms are capable enough to deal with multi-class classification and localization as well as to deal with the objects with multiple occurrences.



Object Detection But real time on Microcontroller



It was only a few years ago that you needed **huge computers** in order to perform this task and now we can actually do it on **microcontrollers**

Different Model

Choose a different model

Did you know? You can customize your model through the Expert view (click on  to switch), or can even bring your own model (in PyTorch, Keras or scikit-learn).

MODEL	AUTHOR	
MobileNetV2 SSD FPN-Lite (320x320 only) <small>OFFICIALLY SUPPORTED</small>	Edge Impulse	Add
A pre-trained object detection model designed to locate up to 10 objects within an image, outputting a bounding box for each object detected. The model is around 3.7MB in size. It supports an RGB input at 320x320px. For other resolutions, use FOMO or an NVIDIA TAO model.		
FOMO (Faster Objects, More Objects) MobileNetV2 0.1 <small>OFFICIALLY SUPPORTED</small>	Edge Impulse	Add
An object detection model based on MobileNetV2 (alpha 0.1) designed to coarsely segment an image into a grid of background vs objects of interest. These models are designed to be <100KB in size and support a grayscale or RGB input at any resolution.		
FOMO (Faster Objects, More Objects) MobileNetV2 0.35 <small>OFFICIALLY SUPPORTED</small>	Edge Impulse	Add
An object detection model based on MobileNetV2 (alpha 0.35) designed to coarsely segment an image into a grid of background vs objects of interest. These models are designed to be <100KB in size and support a grayscale or RGB input at any resolution.		
YOLOv5 for Renesas DRP-AI <small>COMMUNITY</small>	Renesas	Add
Transfer learning model using YOLOv5 v5 branch with yolov5s.pt weights. This block is only compatible with Renesas DRP-AI.		
YOLOv5 <small>COMMUNITY</small>	Community blocks	Add
Transfer learning model based on Ultralytics YOLOv5 using yolov5n.pt weights, supports RGB input at any resolution (square images only).		
YOLOX for TI TDA4VM <small>COMMUNITY</small>	Texas Instruments	Add
TI's EDGEAI YOLOX. https://github.com/TexasInstruments/edgeai-yolox . Outputs ONNX v7 model format both with and without final detect layers using PyTorch 1.7.1. See the implementation https://github.com/edgeimpulse/example-custom-ml-block-ti-yolox/tree/onnx-v7		

Mobilenet Model

Mobilenet is a type of **convolutional neural network** designed for mobile and embedded vision applications. Instead of using standard convolution layers, they are based on a streamlined architecture that uses depthwise separable convolutions.

Using this architecture, we can build lightweight deep neural networks that have low latency for **mobile and embedded devices (example: jetson nano)**.

YOLO Model

YOLO (You Only Look Once) is a **real-time object detection** algorithm developed by **Joseph Redmon and Ali Farhadi** in 2015. It is a single-stage object detector that uses a **convolutional neural network (CNN)** to predict the bounding boxes and class probabilities of objects in input images.

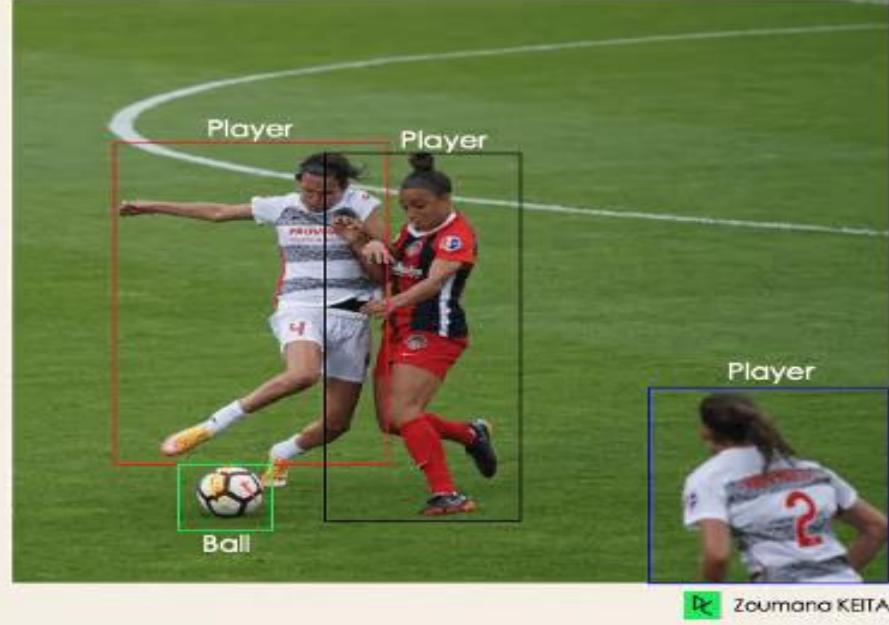
The YOLO algorithm divides the input image into a grid of cells, and for each cell, it predicts the probability of the presence of an object and the bounding box coordinates of the object. It also predicts the class of the object. Unlike two-stage object detectors such as R-CNN and its variants, YOLO processes the entire image in one pass, making it faster and more efficient.

YOLO has been developed in several versions, such as YOLOv1, YOLOv2, YOLOv3, YOLOv4, YOLOv5, YOLOv6, and YOLOv7. Each version has been built on top of the previous version with enhanced features such as improved accuracy, faster processing, and better handling of small objects.

(A) Input image



(B) YOLO Algorithm result

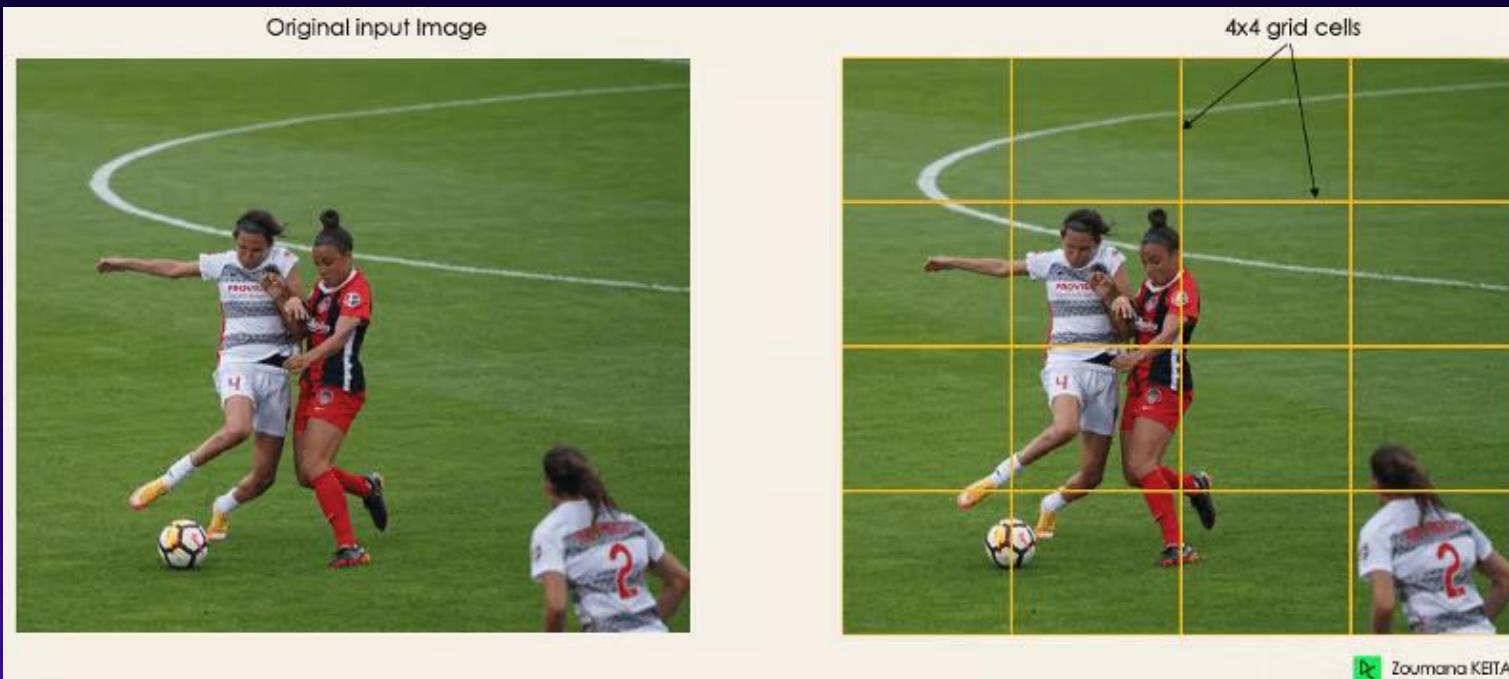


The algorithm works based on the following three approaches:

- Residual blocks
- Bounding box regression
- Intersection Over Unions or IOU for short

Residual blocks

This first step starts by dividing the original image (A) into $N \times N$ grid cells of equal shape, where N , in our case, is 4, as shown in the image on the right. Each cell in the grid is responsible for localizing and predicting the class of the object that it covers, along with the probability/confidence value.



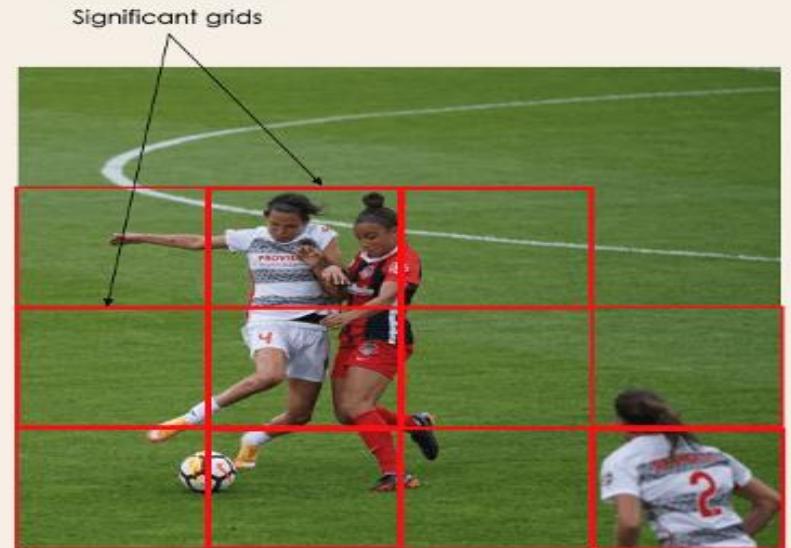
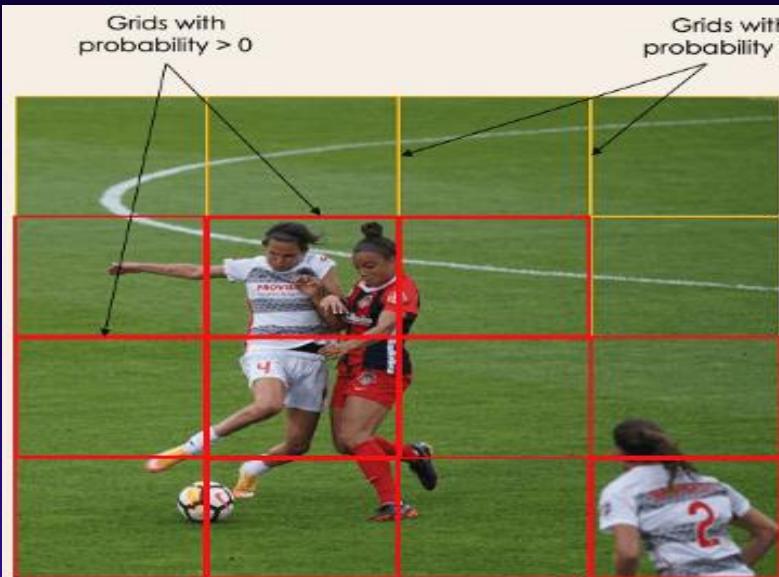
Bounding box regression

The next step is to determine the bounding boxes corresponding to rectangles, highlighting all the objects in the image. We can have as many bounding boxes as there are objects within a given image.

YOLO determines the attributes of these bounding boxes using a single regression module in the following format, where Y is the final vector representation for each bounding box.

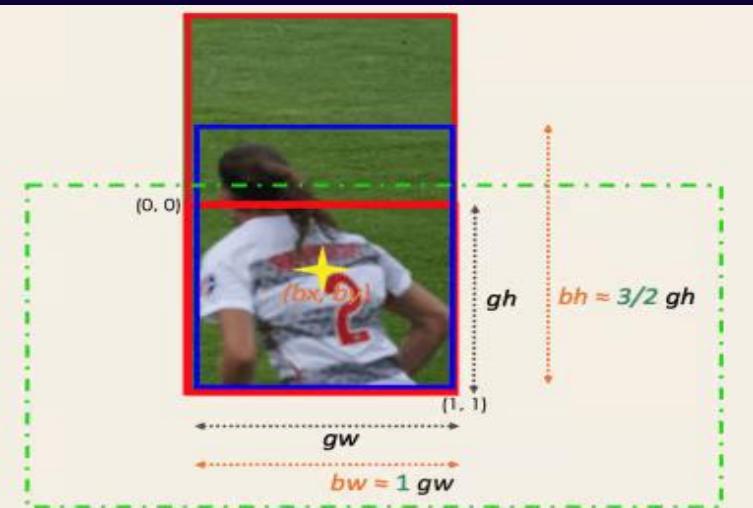
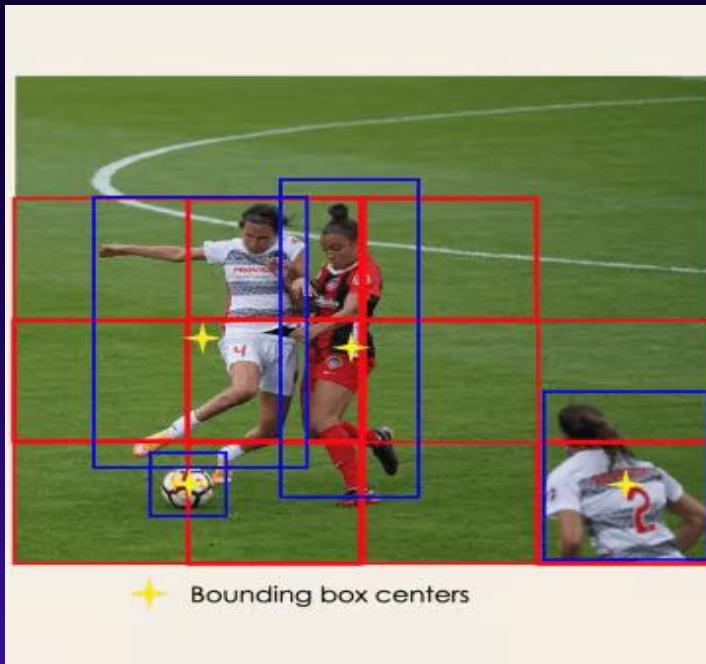
$$Y = [pc, bx, by, bh, bw, c1, c2]$$

- pc corresponds to the probability score of the grid containing an object. For instance, all the grids in red will have a probability score higher than zero. The image on the right is the simplified version since the probability of each yellow cell is zero (insignificant).



- ❑ bx, by are the x and y coordinates of the center of the bounding box with respect to the enveloping grid cell.
- ❑ bh, bw correspond to the height and the width of the bounding box with respect to the enveloping grid cell.

$c1$ and $c2$ correspond to the two classes, Player and Ball. We can have as many classes as your use case requires.



From the previous info we can have for e.g.
 $Y = [1, bx, by, 3/2, 1, c1, c2]$

- First 1 means 100% of object presence

- gh, gw : height & width of the grid
- $0 \leq bx \leq 1$
- $0 \leq by \leq 1$
- bh and bw can be more than 1

Intersection Over Unions or IOU

Most of the time, a single object in an image can have multiple grid box candidates for prediction, even though not all are relevant. The goal of the IOU (a value between 0 and 1) is to discard such grid boxes to only keep those that are relevant. Here is the logic behind it:

The user defines its IOU selection threshold, which can be, for instance, 0.5.

Then, YOLO computes the IOU of each grid cell, which is the Intersection area divided by the Union Area.

Finally, it ignores the prediction of the grid cells having an $\text{IOU} \leq \text{threshold}$ and considers those with an $\text{IOU} > \text{threshold}$.



Applications

Healthcare

Agriculture

Security Surveillance

Self-driving cars

FOMO Model

This algorithm locates the presence of an object in the image and represents it with a **bounding box**. It takes an image as input and outputs the location of the bounding box in the form of **(position, height, and width)**.

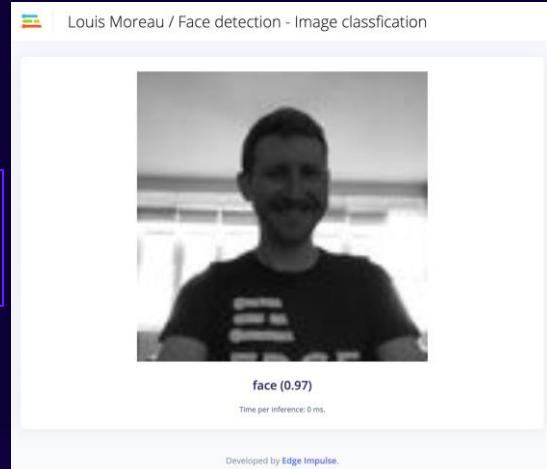
FOMO is a ground-breaking algorithm that brings real-time object detection, tracking and counting to microcontrollers for the first time. FOMO is 30x faster than MobileNet SSD and runs in <200K of RAM.

To give you an idea, we have seen results around 30 fps on the Arduino Nicla Vision (Cortex-M7 MCU) using 245K RAM.

The two most common image processing problems are **image classification** and **object detection**.

Image classification takes an image as an input and outputs what type of object is in the image. This technique works great, even on microcontrollers, as long as we only need to detect a single object in the image.

“Is there a face or not in the image?”



Louis Moreau / Face detection - Image classification

face (0.97)

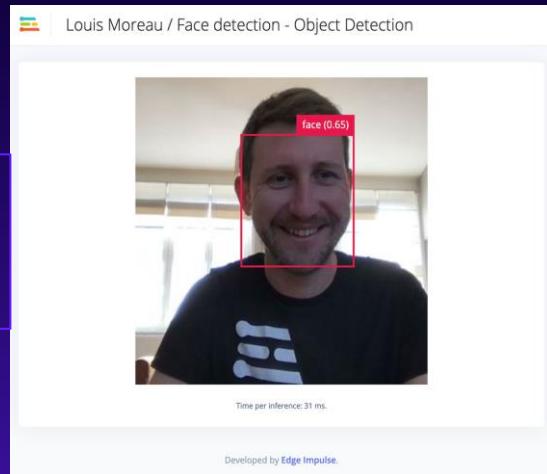
Time per inference: 0 ms.

Developed by Edge Impulse.

This screenshot shows a grayscale portrait of a man. A white rectangular box highlights his face, and the text "face (0.97)" is displayed below it, indicating a high confidence level for face detection. Below the image, the text "Time per inference: 0 ms." and "Developed by Edge Impulse." are visible.

Object detection takes an image and outputs information about the class and number of objects, position, and size in the image.

“Are there faces in the image, where and what size are they?”



Louis Moreau / Face detection - Object Detection

face (0.65)

Time per inference: 31 ms.

Developed by Edge Impulse.

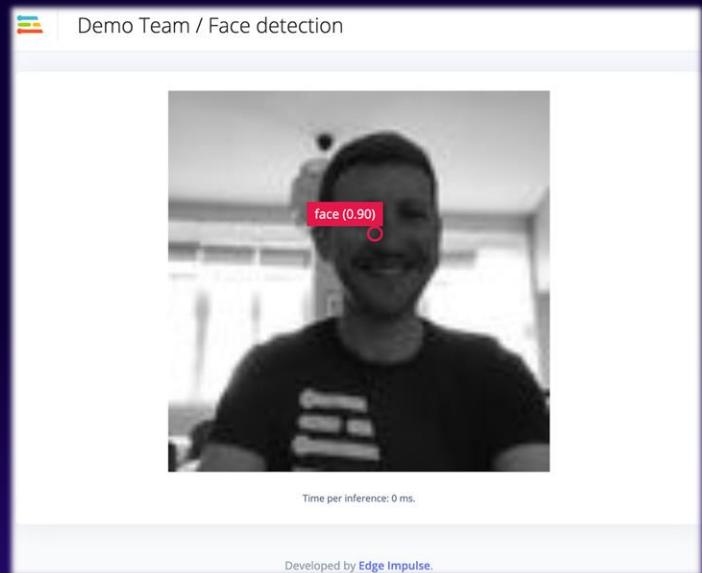
This screenshot shows a grayscale portrait of a smiling man. A red rectangular bounding box surrounds his face, and the text "face (0.65)" is displayed inside the box, indicating a moderate confidence level for face detection. Below the image, the text "Time per inference: 31 ms." and "Developed by Edge Impulse." are visible.

Trained on centroid

The main design decisions for **FOMO** are based on the idea that a lot of object detection problems don't actually need the size of the objects but rather just their location in the frame.

Once we know the location of salient objects we can ask further questions such as "Is an object above/below another?" or "How many of these objects are in view?"

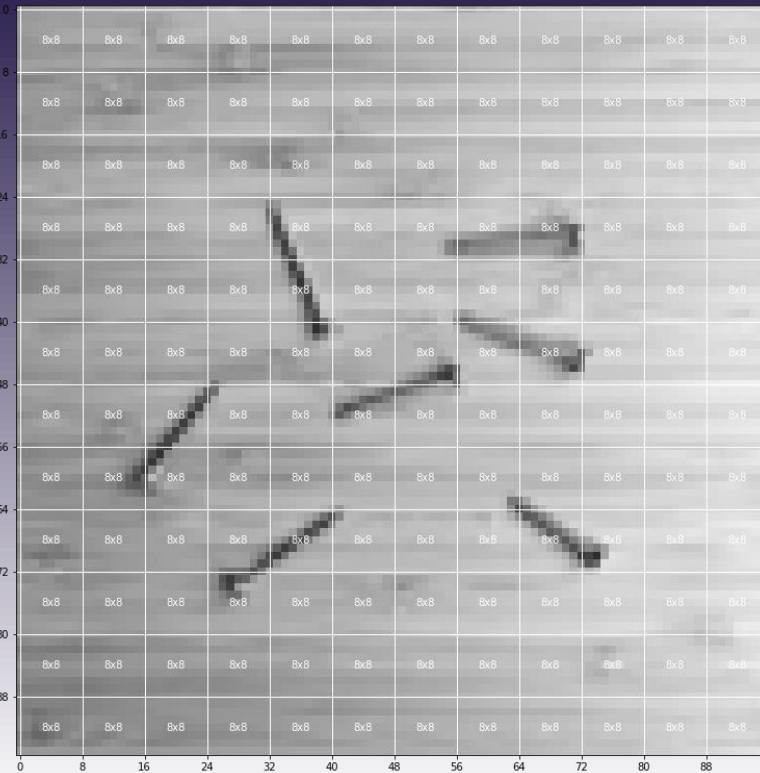
Based on this observation, classical bounding boxes are no longer needed. Instead, a detection based on the centroids of objects is enough.



The question the model is trying to answer is:
"Are there faces in the image, where are they?"

Trained on centroid

when we run the model on an input image, we split this image into a grid and run the equivalent of image classification across all cells in the grid independently in parallel. By default the grid size is **8x8 pixels**, which means for a **96x96 image, the output will be 12x12**. For a 320x320 image, the output will be 40x40.



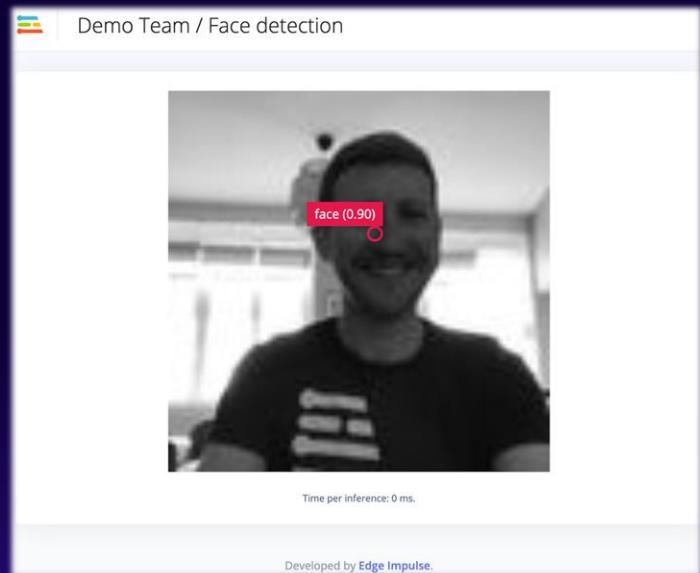
96x96 image split into a 12x12 grid

Trained on centroid

The main design decisions for **FOMO** are based on the idea that a lot of object detection problems don't actually need the size of the objects but rather just their location in the frame.

Once we know the location of salient objects we can ask further questions such as "Is an object above/below another?" or "How many of these objects are in view?"

Based on this observation, classical bounding boxes are no longer needed. Instead, a detection based on the centroids of objects is enough.



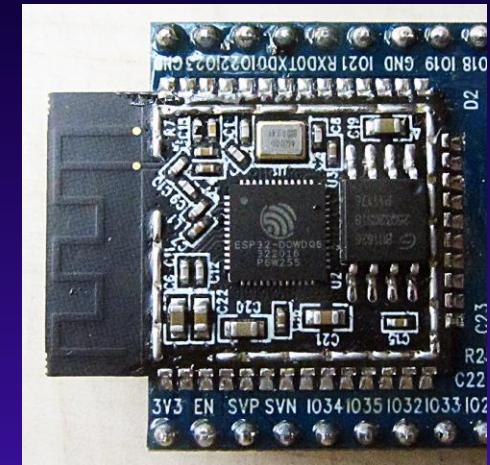
The question the model is trying to answer is:
"Are there faces in the image, where are they?"

Designed to be small and amazingly fast

One of the first goals, when we started to design FOMO, was to run object detection on microcontrollers where flash and RAM are most of the time very limited.

The smallest version of FOMO (96x96 grayscale input, MobileNetV2 0.05 alpha) runs in <100KB RAM and ~10 fps on a Cortex-M4F at 80MHz.

Core	Dual-core 32-bit MCU
Microprocessors	Two Xtensa 32-bit LX6 microprocessors: core 0 and core 1
CPU clock frequency	Adjustable from 80 MHz to 240 MHz
Low-power coprocessor	Yes, can be used instead of the CPU to save power

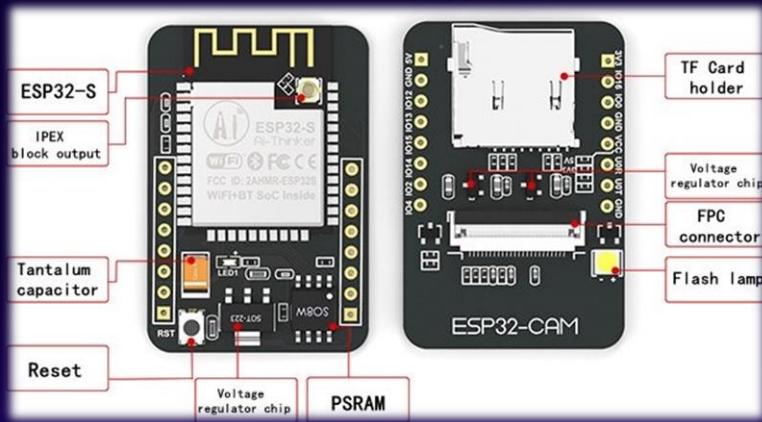


Limitations

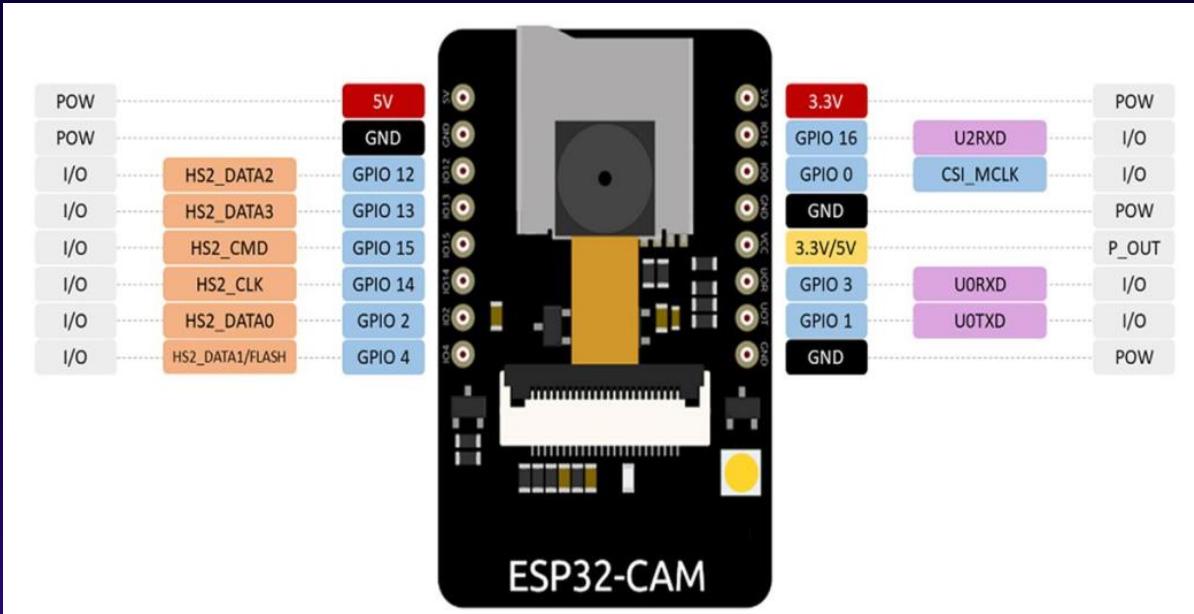
Works better if the objects have a similar size

Objects shouldn't be too close to each other.

ESP32_CAM

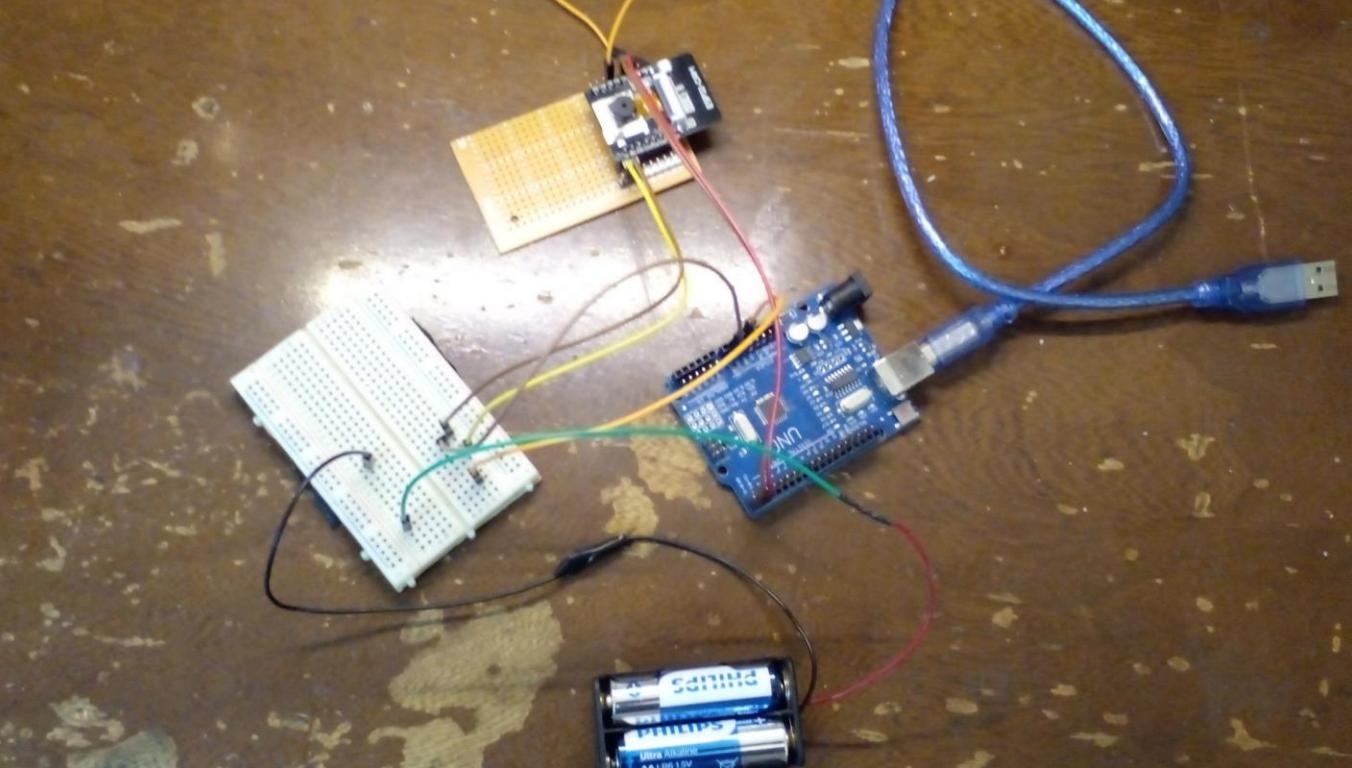
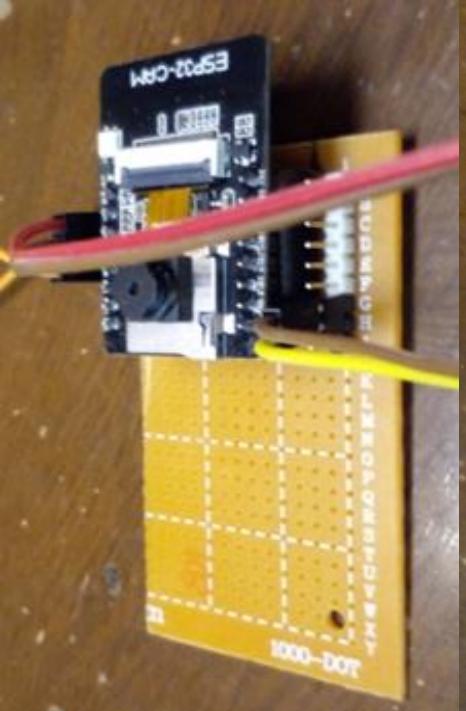


ESP32_CAM



ESP32_CAM

- Product Name: ESP32-CAM.
- WiFi+Bluetooth module: ESP-32S.
- Camera Module: OV2640 2MP.
- Flash Light: LED Built-in on Board.
- Operating Voltage: 3.3/5 Vdc.
- Onboard TF card slot, supports up to 4G TF card for data storage .
- RAM: Internal 512KB + External 4MB PSRAM.
- Power consumption:
 - Flash off: 180mA@5V.
 - Flash on and brightness max: 310mA@5V.
 - Deep-Sleep: as low as 6mA@5V.
 - Modern-Sleep: as low as 20mA@5V.
 - Light-Sleep: as low as 6.7mA@5V
- Dimensions: 40.5mm x 27mm x 4.5mm



Esp32 have the capability of connect to wifi and then stream on a local ip.
You can enter the IP in your browser and watch the stream.

The vey first Picture I took at 3 or 4 PM
after lots of misery



≡ Toggle OV2640 settings

XCLK MHz	20	Set
Resolution	QVGA(320x240)	<input type="button" value="▼"/>
Quality	4	<input type="button" value="▼"/>
Brightness	-2	<input type="button" value="▼"/>
Contrast	-2	<input type="button" value="▼"/>
Saturation	-2	<input type="button" value="▼"/>
Special Effect	No Effect	<input type="button" value="▼"/>
AWB	<input type="checkbox"/>	<input checked="" type="checkbox"/>
AWB Gain	<input type="checkbox"/>	<input checked="" type="checkbox"/>
WB Mode	Auto	<input type="button" value="▼"/>
AEC SENSOR	<input type="checkbox"/>	<input checked="" type="checkbox"/>
AEC DSP	<input type="checkbox"/>	<input checked="" type="checkbox"/>
AE Level	-2	<input type="button" value="▼"/>
AGC	<input type="checkbox"/>	<input checked="" type="checkbox"/>
Gain Ceiling	2x	<input type="button" value="▼"/>
BPC	<input type="checkbox"/>	<input checked="" type="checkbox"/>
WPC	<input type="checkbox"/>	<input checked="" type="checkbox"/>
Raw GMA	<input type="checkbox"/>	<input checked="" type="checkbox"/>
Lens Correction	<input type="checkbox"/>	<input checked="" type="checkbox"/>
H-Mirror	<input type="checkbox"/>	<input checked="" type="checkbox"/>
V-Flip	<input type="checkbox"/>	<input checked="" type="checkbox"/>
DCW (Downsize EN)	<input type="checkbox"/>	<input checked="" type="checkbox"/>
Color Bar	<input type="checkbox"/>	<input checked="" type="checkbox"/>
LED Intensity	0	<input type="button" value="▼"/>
Face Detection	<input type="checkbox"/>	<input checked="" type="checkbox"/>
Face Recognition	<input type="checkbox"/>	<input checked="" type="checkbox"/>

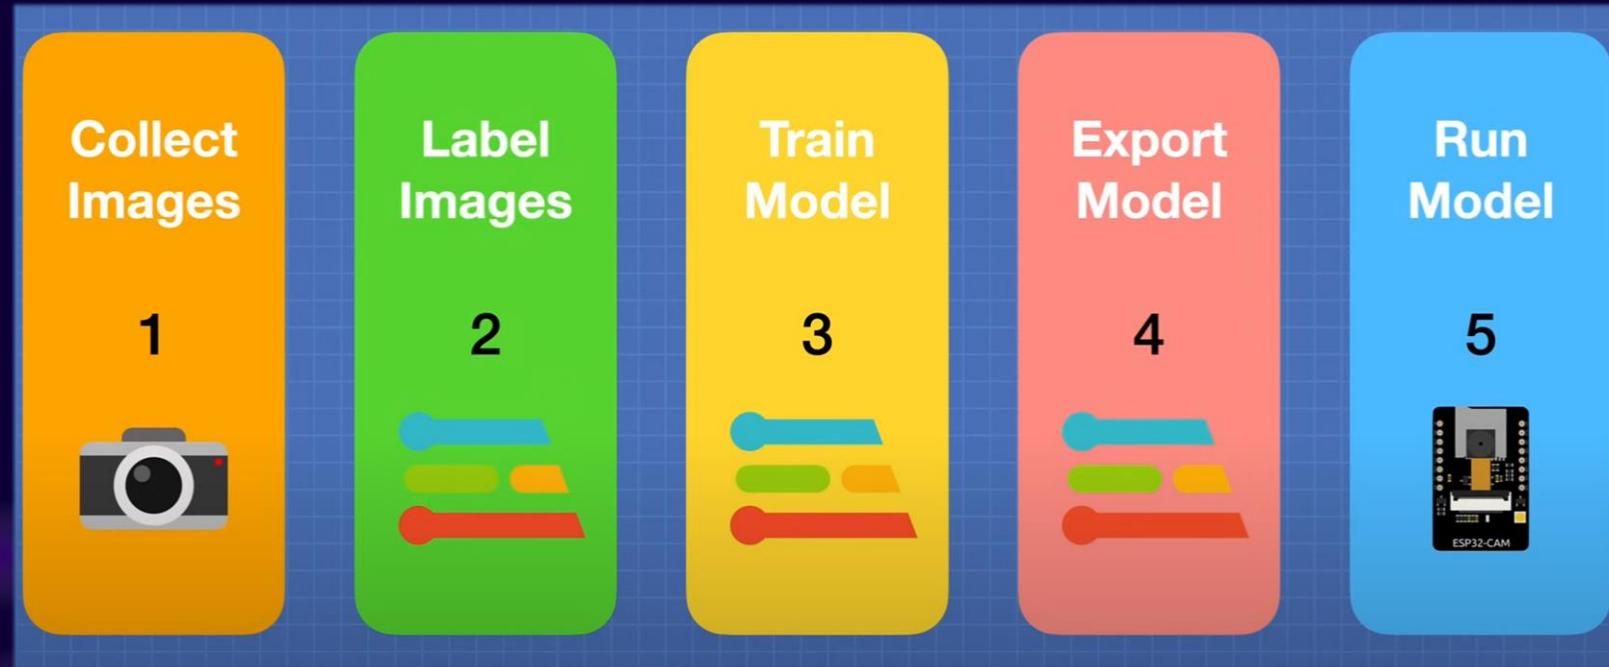


≡ Toggle OV2640 settings

XCLK MHz	20	Set
Resolution	QVGA(320x240)	<input type="button" value="▼"/>
Quality	4	<input type="button" value="▼"/>
Brightness	-2	<input type="button" value="▼"/>
Contrast	-2	<input type="button" value="▼"/>
Saturation	-2	<input type="button" value="▼"/>
Special Effect	No Effect	<input type="button" value="▼"/>
AWB	<input type="checkbox"/>	<input checked="" type="checkbox"/>
AWB Gain	<input type="checkbox"/>	<input checked="" type="checkbox"/>
WB Mode	Auto	<input type="button" value="▼"/>
AEC SENSOR	<input type="checkbox"/>	<input checked="" type="checkbox"/>
AEC DSP	<input type="checkbox"/>	<input checked="" type="checkbox"/>
AE Level	-2	<input type="button" value="▼"/>
AGC	<input type="checkbox"/>	<input checked="" type="checkbox"/>
Gain Ceiling	2x	<input type="button" value="▼"/>
BPC	<input type="checkbox"/>	<input checked="" type="checkbox"/>
WPC	<input type="checkbox"/>	<input checked="" type="checkbox"/>
Raw GMA	<input type="checkbox"/>	<input checked="" type="checkbox"/>
Lens Correction	<input type="checkbox"/>	<input checked="" type="checkbox"/>
H-Mirror	<input type="checkbox"/>	<input checked="" type="checkbox"/>
V-Flip	<input type="checkbox"/>	<input checked="" type="checkbox"/>
DCW (Downsize EN)	<input type="checkbox"/>	<input checked="" type="checkbox"/>
Color Bar	<input type="checkbox"/>	<input checked="" type="checkbox"/>
LED Intensity	0	<input type="button" value="▼"/>
Face Detection	<input type="checkbox"/>	<input checked="" type="checkbox"/>
Face Recognition	<input type="checkbox"/>	<input checked="" type="checkbox"/>



Main Steps for Object detection using ESP32-CAM



Dashboard

Devices

Data acquisition

Experiments

EON Tuner

Impulse #1

Create impulse

Image

Object detection

Upgrade Plan

Get access to higher job limits, collaborators and a full commercial license.

[View plans](#)[Dataset](#) [Data sources](#) | [Labeling queue \(0\)](#) [AI labeling](#)

DATA COLLECTED

181 items



TRAIN / TEST SPLIT

75% / 25%



Dataset

Training (135) Test (46)

SAMPLE NAME

LABELS

ADDED

Scissors_1734177027309

Scissors

Dec 15 2024, 2...

⋮

Scissors_1734177018812

Scissors

Dec 15 2024, 2...

⋮

Scissors_1734177018575

Scissors

Dec 15 2024, 2...

⋮

Scissors_1734177018307

Scissors

Dec 15 2024, 2...

⋮

Scissors_1734177010809

Scissors

Dec 15 2024, 2...

⋮

Scissors_1734177003809

Scissors

Dec 15 2024, 2...

⋮

Scissors_1734177010566

Scissors

Dec 15 2024, 2...

⋮

Collect data

[Connect a device](#) to start building your dataset.

RAW DATA

Scissors_1734177027309

Metadata

?

+

Impulse #1



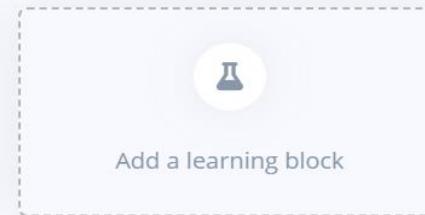
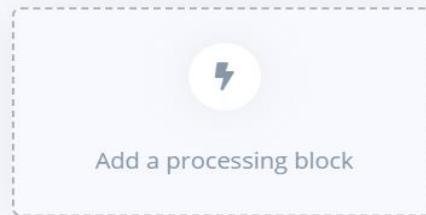
An impulse takes raw data, uses signal processing to extract features, and then uses a learning block to classify new data.

Image data

Input axes
image

Image width **Image height**

Resize mode



Output features



⚡ Add a processing block

Did you know? You can bring your own DSP code.

DESCRIPTION	AUTHOR	RECOMMENDED	
Image OFFICIALLY SUPPORTED Preprocess and normalize image data, and optionally reduce the color depth.	Edge Impulse	★	Add
Flatten OFFICIALLY SUPPORTED Flatten an axis into a single value, useful for slow-moving averages like temperature data, in combination with other blocks.	Edge Impulse		Add
Audio (MFCC) OFFICIALLY SUPPORTED Extracts features from audio signals using Mel Frequency Cepstral Coefficients, great for human voice.	Edge Impulse		Add
Audio (MFE) OFFICIALLY SUPPORTED Extracts a spectrogram from audio signals using Mel-filterbank energy features, great for non-voice audio.	Edge Impulse		Add
Spectral Analysis OFFICIALLY SUPPORTED Great for analyzing repetitive motion, such as data from accelerometers. Extracts the frequency and power characteristics of a signal over time.	Edge Impulse		Add
Spectrogram OFFICIALLY SUPPORTED Extracts a spectrogram from audio or sensor data, great for non-voice audio or data with continuous frequencies.	Edge Impulse		Add
DESCRIPTION	AUTHOR	RECOMMENDED	
Object Detection (Images) OFFICIALLY SUPPORTED Fine tune a pre-trained object detection model on your data. Good performance even with relatively small image datasets.	Edge Impulse	★	Add
Object Detection (Images) - BrainChip Akida™ OFFICIALLY SUPPORTED Fine tune a pre-trained object detection model on your data. Good performance even with relatively small image datasets. Only works with BrainChip AKD1000 MINI PCIe board.	BrainChip	★	Add
Classification OFFICIALLY SUPPORTED Learns patterns from data, and can apply these to new data. Great for categorizing movement or recognizing audio.	Edge Impulse		Add
Transfer Learning (Images) OFFICIALLY SUPPORTED Fine tune a pre-trained image classification model on your data. Good performance even with relatively small image datasets.	Edge Impulse		Add
Regression OFFICIALLY SUPPORTED Learns patterns from data, and can apply these to new data. Great for predicting numeric continuous values.	Edge Impulse		Add
Transfer Learning (Keyword Spotting) OFFICIALLY SUPPORTED Fine tune a pre-trained keyword spotting model on your data. Good performance even with relatively small keyword datasets.	Edge Impulse		Add

Parameters

Generate features

Raw data

Show: All labels

Scissors 1734177027309 (Scissors) ▾



Raw features

0x667268, 0x667268, 0x667268, 0x667268, 0x667268, 0x667268, 0x68746a, 0x69756b, 0x69756b, 0x6b776d, 0xed796f, 0xd796f, 0xd7...

Parameters

Image

Color depth ?

Grayscale

Save parameter

DSP result

Image



Processed features

0.4285, 0.4285, 0.4285, 0.4285, 0.4285, 0.4285, 0.4364, 0.4403, 0.4403, 0.4481, 0.4560, 0.4560, 0.4560, 0.4481, 0.4481, 0.4481,

On-device performance ⓘ



PROCESSING TIME
15 ms.



PEAK RAM USAGE

Parameters

Generate features

Training set

Data in training set 135 items

Classes 2 (Apple, Scissors)

Generate features

Feature generation output

0 (0)

Job started at 17 Dec 2024 19:06:48

Reducing dimensions for visualizations...

UMAP(verbose=True)

Tue Dec 17 19:07:00 2024 Construct fuzzy simplicial set

Tue Dec 17 19:07:00 2024 Finding Nearest Neighbors

Tue Dec 17 19:07:02 2024 Finished Nearest Neighbor Search

Tue Dec 17 19:07:04 2024 Construct embedding

Epochs completed: 100% 500/500 [00:02<00:00, 222.86it/s]

Tue Dec 17 19:07:06 2024 Finished embedding

Writing output files...

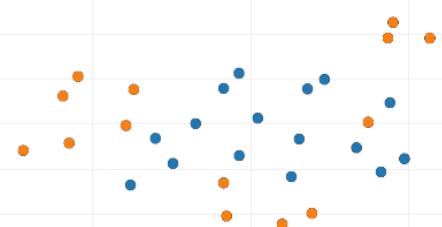
Writing output files OK

Reducing dimensions for visualizations OK (took 801ms.)

Job completed (success)

Feature explorer

Apple
Scissors



On-device performance



PROCESSING TIME
15 ms.



PEAK RAM USAGE
4 KB

Neural Network settings

Training settings

Number of training cycles 

Use learned optimizer 



Learning rate 

Training processor 

Data augmentation 



Advanced training settings

Neural network architecture

Input layer (9,216 features)



FOMO (Faster Objects, More Objects) MobileNetV2 0.1

Choose a different model

Output layer (2 classes)

Save & train

Training output

Model

Model version:  

Last training performance (validation set)



F1 SCORE 

83.3%

Confusion matrix (validation set)

	BACKGROUND	APPLE	SCISSORS
BACKGROUND	100%	0%	0%
APPLE	25%	75%	0%
SCISSORS	33.3%	0%	66.7%
F1 SCORE	1.00	0.86	0.80

Metrics (validation set)

METRIC	VALUE
Precision (non-background) 	1.00
Recall (non-background) 	0.71
F1 Score (non-background) 	0.83

On-device performance 

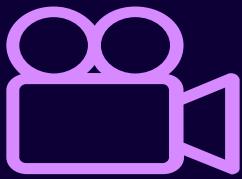
Engine:  EON™ Compiler 

 INFERENCE TIME
937 ms.

 PEAK RAM USAGE
235.5K

 FLASH USAGE
65.0K





Video

Resources

www.datacamp.com

www.studio.edgeimpulse.com

www.edgeimpulse.com

www.youtube.com

www.docs.edgeimpulse.com

www.geeksforgeeks.org

www.geeksforgeeks.org

www.geeksforgeeks.org

www.geeksforgeeks.org